

CONSTRUCTOR UNIVERSITY BREMEN

Bachelor of Science
Computer Science

Andrei Kozyrev

Equality saturation for solving equalities of relational expressions

Bachelor's Thesis

Scientific supervisor:
professor Anton Podkopaev

Reviewer:

Bremen
2023

Abstract

Modern CPUs are being developed exceptionally fast, and the number of cores is increasing rapidly. This has led to the development of multithreading, which is a technique that allows for the execution of multiple threads on a single CPU. Memory models are a fundamental aspect of multithreading and describe how memory is ordered at runtime in relation to source code. Currently, existing memory models are unsatisfactory and there is a need for new models that can be rigorously proven. In order to achieve this, formal verification using the Coq proof assistant is utilized, which enables automated proof checking and ensures the accuracy of results. Specialists in weak memory are continuously improving the results in this domain.

One of the big and common problems in weak memory is the proof of equivalence of several memory models. Memory models are represented as expressions over relational language.

This thesis focuses on the automation of proving equalities over relational expressions in Coq. We are utilizing the techniques of equality saturation and E-graph data structure to generate proof of equivalence for a given pair of terms. By automating these proofs, we can greatly increase the efficiency and accuracy of the proof process in weak memory.

Acknowledgements

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Contents

Abstract	3
Acknowledgements	5
1. Introduction	8
1.1. Approach	10
1.2. Core implementation components	11
2. Objectives	13
3. Related Work	14
3.1. Coq	14
3.1.1. Coq Overview	14
3.1.2. Coq Proof mode	14
3.1.3. Coq’s Significance	16
3.1.4. Proof Automation	16
3.2. Alternative solutions to solving relational equations	17
3.3. E-graphs and Equality Saturation	18
4. Implementation	20
4.1. Vernacular Commands	21
4.2. OCaml plugin	21
4.3. Using Egg	24
References	25

1 Introduction

Memory model is an abstraction, invented to describe the behavior of a program in a multithreaded environment. Memory models are used to study the behavior of concurrent executions on various computing systems. Memory models consist of a set of rules, which dictate how memory operations interact with each other.

One of the ways to define a memory model is to use declarative semantics, where memory model could be represented as a graph, where nodes represent memory operations and edges represent order on these operations. Given a concrete execution of a program on a given architecture, we can construct a graph from a Happens Before Relation over pairs of Memory Operations. This graph would represent the execution of the program.

The most traditional and conservative memory model is Sequential Consistency. Sequential Consistency states that all memory operations happening in a single thread are related by the Program Order.

“The result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each processor appear in the order specified by its program.”
(Leslie Lamport (1979))

Memory models that are more relaxed than Sequential Consistency are called weak memory models. While Weak Memory models give compilers and processors more flexibility to optimize, they introduce more problems to the programmer.

If we reason about weak memory models in a declarative way, we use relations to describe models and facts about them. Proofs of various model properties also require using relational expressions. Proofs of such propositions are typically massive and very error prone. There were several cases of incorrect result in submitted and published papers. Batty et al. [4] suggested an incorrect fix for the semantics of SC calls in C++, which was

later documented by Lahav et al. [8]. Moreover Pichon-Pharabod et al. in 2016 suggested an incorrect proof of compilation in their paper [13].

Given that writing and maintaining proofs about Weak Memory on paper is very difficult, it is considered good practice to write them in Coq [5], which is a proof assistant system. Coq helps to grant the correctness of the proof process. In Coq, proofs are expressed as formal mathematical objects, and the correctness of the proof is ensured by checking that the proof is consistent with the axioms and rules of logic. Coq is a tool that facilitates the creation of mechanically-verifiable proofs, in contrast to paper-based proofs which are prone to errors.

Weak memory proofs, even with a use of Coq, tend to be huge and convoluted. This is why in this thesis, we are focusing on automating a specific part of weak memory proofs, namely the proofs of equivalences between several memory models. As already mentioned, we can define memory models as axioms over relational language. Consider an example of a proposition about relations. Let us denote $\mathbf{r}, \mathbf{r}' \subset \mathbf{A} \times \mathbf{A}$ — two relations over a given set \mathbf{A} . Now we may consider common operations in relational language, e.g. transitive closure (\mathbf{r}^*), reflexive closure ($\mathbf{r}^?$), or composition ($\mathbf{r} ;; \mathbf{r}'$). An example of a proposition we want to prove may look like this:

$$(\mathbf{r}^* ;; \mathbf{r}^?) ;; \mathbf{r}^? \equiv \mathbf{r}^*$$

If we were to prove this statement in Coq by hand, we would have to consequently apply multiple theorems to rewrite both sides of the equivalence relation until they are syntactically equal. In this particular example, we would have to apply theorem `rt_cr` twice:

```
rt_cr : forall (A : Type) (r : relation A), r* ;; r? = r*
```

After `rt_cr` being applied to the left hand side (lhs) two times, the two sides become syntactically equal. `rt_cr` and other theorems that help to reason about relations are provided by the Hahn library [1].

The general problem discussed in this thesis is as follows: given a rewriting system, i.e. a set of theorems, and an equivalence relation between two expressions, we want to find a sequence of rewrites that can be applied to both sides of the relation to make them syntactically equal. This is a wider problem than reasoning about relational expressions, thereafter the solution we propose could be adapted and used to solve other problems that involve derivability in two-sided associative calculus over given language in Coq. This thesis, in turn, focuses on applying the proposed technique to automate weak memory, where the rewriting system we use is the Hahn library.

As a basis for various proofs in the area of weak memory, the weakmemory¹ organization on GitHub, specifically the imm² repository, was utilized. The work on the thesis involves analyzing present proofs and their possible patterns to automate the process and attempt to shorten them using the developed tool.

1.1 Approach

Our approach to solve equivalences is based on the technique called *equality saturation* which utilizes the data structure, called an *e-graph*. E-graph stores an equivalence relation over terms of some language and allows to store potentially exponential amount of terms in a compact way. E-graph is a set of equivalence classes (*e-classes*) and e-class is a set of *e-nodes*, which represent equivalent terms in the language. Whilst e-nodes are function symbols, associated with a list of e-classes.

Equality saturation is a process of iteratively applying a set of rewrite rules to the e-graph until rewrites bring no more new information to the graph. That point is called saturation. The saturated e-graph represents a set of all possible terms equivalent to the origin, that can be obtained by applying the rewrite rules to the initial term.

¹<https://github.com/weakmemory>

²<https://github.com/weakmemory/imm>

In a saturated e-graph it becomes algorithmically easy to check if two terms are equivalent and, moreover, to find a proof of their equivalence. Having a sequence of rewrites, we can apply them within the Coq proof view. Just as we would do by hand, applying the theorems one by one.

1.2 Core implementation components

We have chosen fast and lightweight **egg** [15] rust library as an equality saturation engine for our project. **Egg** is a new and ambitious project, that has already proven its worth and was used in tools like *Ruler* [10].

In Coq you can add new functionality by writing a plugin. Plugins are written in OCaml language. We have set the communication between OCaml and Rust using the **OCaml-rs** [2] foreign function interface (FFI). To prove theorems in Coq, you use proof mode, which is activated when you start a proof, for example, by using the **Theorem** command. The proof state contains one or more unproven goals. Each goal consists of a conclusion, which is the statement to be proven, and a local context. When the proof is complete, you exit proof mode by using the **Qed** command. During proof mode, you use so called *tactics*, to gradually transform the proof. Tactics specify how to transform the proof state of an incomplete proof and are used to eventually generate a complete proof.

Commands in Coq are similar to tactics, but accessible from outside of proof mode, e.g. command **Print**, which prints the given term to console. A Coq plugin typically consists of a set of commands and tactics, that are used from Coq as a front-end to the plugin. We have defined multiple tools inside our plugin, that are used to interact with the **egg** library:

- **Cegg solve** — a tactic that simplifies the lhs of the equation. It takes the conclusion of the goal, retrieves the lhs of the equivalence, parses it into an s-expression³ and passes it to **egg**. **Egg** builds an e-graph E for it and extracts the “best” term t from E . A sequence of rewrites

³S-expressions represent tree-like data

to achieve t is passed back to OCaml and is applied to the lhs of the equivalence inside Coq proof mode.

- `Cegg solve eq` — a tactic that tries to prove the equivalence between the lhs and rhs of the equation, using `egg`.
- `Cegg config` — a command, that allows to configure the ruleset for `egg`. It takes a user-defined list of rewrite rules and caches it for the later use in `Cegg solve` and `Cegg solve eq`.

The source code and documentation for the thesis is available at:

<https://github.com/K-dizzled/relations-via-egg>

2 Objectives

This thesis aims to automate solving equivalences over relational expressions in Coq. The goal is to manage to use the **egg** library and automate the proof process. The main objectives are:

- Develop a framework for communication between Rust and Coq interactive proofs.
- Make a rule-parameterized algorithm on top of the **egg** library for producing a series of rewrites, which prove relational equivalence. Given two expressions and a rewriting system, there may be multiple approaches to prove the equivalence.
- Experiment on existing proofs and used lemmas to come up with a usable and efficient rule set. Firstly, to analyse existing weak memory proofs and provide users a useful interface. Secondly, as big rule sets result in huge e-graphs, that take long to be built, we aim to research on the efficient and substantive rule set, that would be small enough to be used to build e-graphs in practice.

3 Related Work

First part of this chapter focuses on introducing the reader to Coq proof assistant, its core concepts and significance. Then other solutions to the problem are discussed in Section 3.2. Finally, we focus on the details of our particular approach and cover details regarding equality saturation and the `egg` library in Section 3.3.

3.1 Coq

This section aims to provide an overview of Coq and its significance in the field of theorem proving. We will outline the advantages of using Coq compared to other proof assistants and methods. Additionally, we will delve deeper into the Coq proving process.

3.1.1 Coq Overview

Coq is a formal proof management system. It provides a language called Gallina, which is used to define mathematical objects and write formal proofs. The formalism behind Coq is the Calculus of Inductive Constructions (CIC) [12]. In CIC types are used to ensure the correctness of proofs. Each theorem is essentially a type, and its proof is a value inhabiting that type. It might also be useful to think of proofs as functions from hypothesis to conclusion. For example, assume we have a context Γ , a proposition φ and we want to prove $\Gamma \vdash \varphi$. That would mean that the proof we are looking for is a mapping, which for any argument of type Γ constructs a value of type φ .

3.1.2 Coq Proof mode

As mentioned earlier, proofs can be viewed as functions and this is one of the ways to make a proof in Coq. We will provide an example for a better understanding. Consider the following theorem: given value of type A and a function $f : A \rightarrow B$, we can construct a value of type B .

```

Definition test (A B : Prop) :
  A -> (A -> B) -> B.

```

To prove such theorem we need to apply the value a of type A to the function f and conclude the proof:

```

Definition test (A B : Prop) :
  A -> (A -> B) -> B
:= fun (a : A) (f : A -> B) => f a.

```

Constrating proof terms by hand may be useful to learn Coq, but it is very inconvenient to write bigger proofs in such manner. In a bigger proof, as like as in paper proofs, you want to iteratively modify the environment and step by step achieve the goal. Coq enters proof mode when you begin a proof, such as with the **Theorem** command:

```

Theorem test (A B : Prop) :
  A -> (A -> B) -> B.
Proof.

```

When you enter a proof mode, you are able to always see current unfinished goals and an up-to-date hypotheses:

```

A, B : Prop
=====
A -> (A -> B) -> B

```

Now we can proof the theorem using tactics. Tactics implement backwards reasoning, so when tactic is applied, all hypotheses it uses to modify the goal are added to the context.

Firstly, we call an **intros** tactic, which introduces all propositions on the left side of the implication as assumptions:

```

A, B : Prop
H : A
H0 : A -> B
=====
B

```

We have a hypothesis H of type A , but we need B . We call an **apply**

H0 tactic, which proves us B , but adds A as a new goal. Now if we call `apply H`, we will conclude the proof.

```
Lemma test' (A B : Prop) :  
  A -> (A -> B) -> B.  
Proof.  
  intros.  
  apply H0.  
  apply H.  
Qed.
```

3.1.3 Coq's Significance

There are other Proof Assistants continuing to appear, but Coq has been developed and improved since 1989. The heart of a proof assistant is its kernel, which is the primary source of reliability of the tool. Unfortunately, regarding of the kernel's size, as any other software, it has issues. However, the amount of work done by the Coq community to recheck and validate the kernel, makes Coq the most trustworthy proof assistant available.

Moreover, Coq's huge strength is the amount of impressive results, obtained using it. The most famous result, achieved using Coq, is the Four Color Theorem proof. The Four Color Theorem states that any map can be colored using only four colors, so that no two adjacent regions are colored the same. There were several paper proofs, that were all proven to be incorrect after a while. A computer-assisted proof was proposed by Kenneth Appel and Wolfgang Haken in 1976 [3]. The proof reduced the problem to the analysis of a smaller number of options and their enumeration by a computer for many hours. Consequently, the mathematical community was not inclined to trust the proof. The first formal checked proof was proposed by George Gonthier et al. in 2005 [6], which showed the community that Coq is ready for huge and complex proofs.

3.1.4 Proof Automation

This thesis focuses on automating the proofs, whilst a question may arise whether such approach makes proofs less clear to read. There are several

proof styles in Coq. A framework called **SSReflect**⁴ exists. Conceptually, **SSReflect** differs from ordinary **Tactics**⁵ in that proofs are written with almost no automation, and the tactics language is much more expressive. Even though **SSReflect** has less automation, proofs anyway tend to be confusing. From the other hand, when classical Coq **Tactics** are used, proof is usually broken down into a huge number of lemmas and sub-claims, definitions of which make the idea clear. Lemmas that are deep down inside the proof can typically be automated without sacrificing the readability, because their proofs are self-evident.

3.2 Alternative solutions to solving relational equations

In 2023 Kokologiannakis et al. presented the tool called **Kater** [7]. **Kater** is a tool that allows to automatically answer memory-model questions. **Kater** is a useful, but standalone tool. Firstly, that means we must believe in its correctness, whereas any functionality integrated into Coq is protected by the Coq’s typechecker from producing incorrect results. Secondly, proofs in weak memory are often more diverse than just reasoning about relational expressions. For instance, we may want to prove correctness of the compilation scheme, e.g. from C to Assembly language. In such cases, the semantics of the assembly code must be related to the original operational semantics of the instructions in our language. While relations are still relevant, they are only a small part of the problem. The main focus is on proving properties about the compilation process. Therefore, **Kater** is suitable only for a limited set of problems, where relations are the primary concern, and the proof process does not involve other complex components. In contrast, Coq provides a comprehensive framework for formal reasoning that allows us to tackle a wide range of problems.

⁴<https://coq.inria.fr/refman/proof-engine/ssreflect-proof-language.html>

⁵<https://coq.inria.fr/refman/proof-engine/tactics.html#tactics>

3.3 E-graphs and Equality Saturation

E-graphs are a generalization of a set-union data structure [14]. E-graphs were first introduced by Greg Nelson and Derek C. Oppen in 1980 [11]. Since then, e-graphs have been used for a successful mechanical theorem proving and program verification, e.g. in Stanford Pascal Verifier at 1979 [9].

We will now formally define an e-graph. Let Σ be a set of function symbols. Let Σ_n be a subset of Σ consisting of symbols of arity n . Let Id be a set of unique identifiers $\text{id}_1, \text{id}_2, \dots, \text{id}_k$, called the e-class **Ids**. Then **e-node** is a function symbol $f \in \Sigma$ and a list of n e-class Ids. E-node is denoted $f(\text{id}_1, \text{id}_2, \dots, \text{id}_n)$. E-class is a set of e-nodes. E-graph is a Disjoint Set Union data structure over e-class Ids. Here we give an example of a how a simple e-graph is build. Consider such a simple language:

```
enum SimpleLanguage {
    Num(i32),
    '+' = Add([Id; 2]),
    '*' = Mul([Id; 2]),
    '/' = Div([Id; 2]),
    '<<' = Bitwise([Id; 2]),
    Symbol(Symbol),
}
```

It is denoted as if we were to use it in **egg** library. Elements of the enum are function symbols with a name and a given arity. Notation explained:

$$\underbrace{''+''}_{\textcircled{1}} = \underbrace{\text{Add}}_{\textcircled{2}}(\underbrace{[\text{Id}; 2]}_{\textcircled{3}})$$

- ① — A string literal to automatically generate a parser for the language
- ② — A unique identifier of a function symbol

③ — Arity of the function symbol

Let us now consider the following arithmetic expression: $(a \times 2) / 2$. The resulting e-graph is shown in Figure 1. E-classes are shown in dashed boxes and e-nodes are circles with their function symbols. Edges represent the parent-child relationship between e-nodes. In Figure 1, no rewrite rules were yet added, so each e-node is in its own e-class.




Figure 1




Figure 2




Figure 3




Figure 4

When the e-graph is built we can run an equality sturation algorithm. Consider a rewriting system S , containing rules of form $\mathbf{lhs} \longrightarrow \mathbf{rhs}$, where \mathbf{lhs} and \mathbf{rhs} are expression in our language. On a high level algorithm does the following:

1. Iterate over rules $r \in S$. For each rule r :
 - (a) Try to match the pattern from r 's \mathbf{lhs} with e-nodes present in the graph.
 - (b) If a match was found: if \mathbf{rhs} was not present in the graph, add it. Then merge \mathbf{lhs} and \mathbf{rhs} e-classes.
2. Loop step 1 until each new iteration introduces new equivalences into the graph. After the process is finished, graph is called saturated.

Figures 1–4 illustrate how the algorithm proceeds. In Figure 2 the rule $x \times 2 \longrightarrow x \ll 1$ was applied. In Figure 3 the rule $(x \times y)/z \longrightarrow x \times (y/z)$ was applied. In Figure 4 rules $x/x \longrightarrow 1$ and $1 \times x \longrightarrow x$ were applied.

4 Implementation

Let us summarize the technical problem and the main components of the developed system. Given a proposition about relations in Coq, we want to prove it using equality saturation, performed by the `egg` library in Rust.

A Coq plugin is a tool with external functionality, added to Coq. There are two main approaches of writing Coq plugins:

- `Ltac`⁶ and `Ltac2`⁷ are languages that help to write simple plugins, combining basic combinations of tactics and actions into a single tactic. It is useful, but not powerful enough to write complex plugins.
- The most traditional way of building new complex tactics is to write a Coq plugin in OCaml.

Coq’s compiler is written in OCaml, so plugins written in OCaml allow to extend Coq’s grammar, along with adding complex logic to the tactic, e.g. using FFI to call external libraries. To have a better understanding of how all the components interact with each other in our work, we provide a diagram in Figure 5.

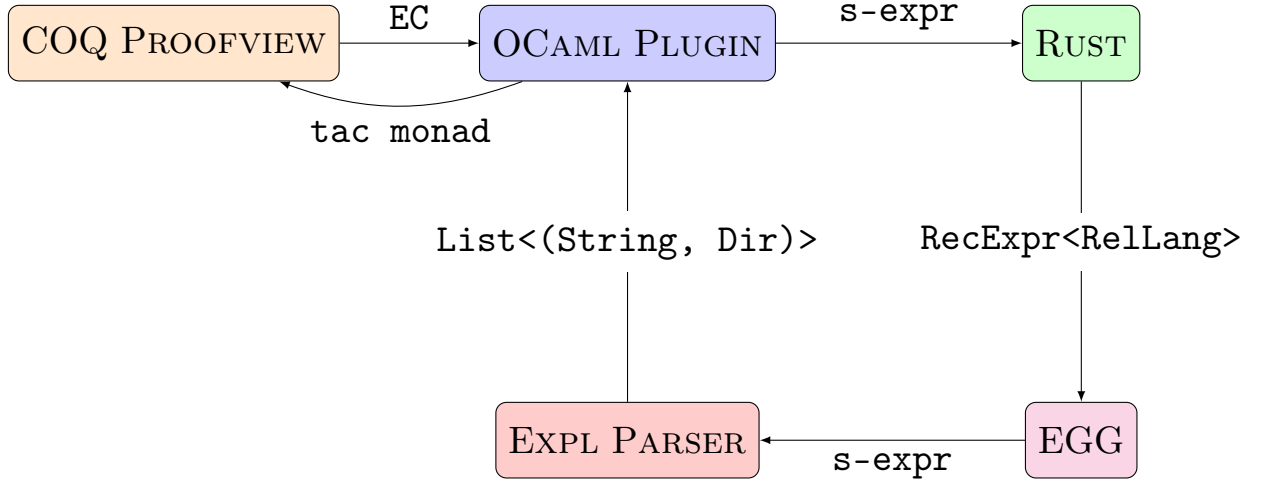


Figure 5: Components of the system

⁶<https://coq.inria.fr/refman/proof-engine/ltac.html#ltac>

⁷<https://coq.inria.fr/refman/proof-engine/ltac2.html#ltac2>

Data types over each arrow from Figure 5 will be discussed in detail in the following chapter.

4.1 Vernacular Commands

To extend Coq’s grammar with a new tactic or command, one should write a `.mlg` file, where the tactic’s syntax is defined. Consider an example:

```
DECLARE PLUGIN "coq-via-egg-plugin.plugin"

VERBAC COMMAND EXTEND cegg_config CLASSIFIED AS QUERY
| [ "Cegg" "config" reference(r) ] -> { ... }
END

TACTIC EXTEND cegg_solve
| [ "Cegg" "solve" ] -> { ... } (* Paring and interpretation rule *)
END
```

In listing 4.1, we define a command called `cegg_config` and a tactic called `cegg_solve`. A command is marked as `QUERY`, meaning it is a *pure* function. Otherwise it would have been marked as `SIDEFF`.

Definition 4.1 (Pure function) *A pure function is a function where the return value is only determined by its input values, without observable side effects.*

After the `|` symbol, we define the parsing rule and the the interpretation rule, separated by `->`. The parsing rule itself is a set of terminals that are matched against a string of tokens. The interpretation rule is a function that is called when the parsing rule is matched. More on the `.mlg` file format can be found in the Coq’s plugin guide ⁸.

4.2 OCaml plugin

When our tactic is called from inside the Coq proof, firstly, we need to extract the goal. Consider an example:

⁸https://github.com/coq/coq/blob/master/doc/plugin_tutorial/tuto2/src/g_tuto2.mlg

```

Lemma example (r : relation A) :
  r* ;; r? ≡ r*.
Proof.
  Cegg solve eq.

```

When we enter the interactive proof process, we see the following proof state:

```

A : Type
r : relation A
=====
r* ;; r? ≡ r*

```

Current hypotheses are located above the line and the conclusion of the goal — below. To interact with Coq from OCaml, we use the Coq-Api⁹, which provides a widest functionality. To extract that information about the goal in OCaml we first enter the Proofview monad.

Definition 4.2 (Monad) *Monads serve as a representation of computations. Consider a computation to be similar to a function that transforms an input to an output, but with an additional component. This component represents the effect that the function has as a consequence of being executed.*

In the following function `t` denotes the goal. `enter` applies the goal-dependent tactic, with `(t -> unit tactic)` type in each goal independently.

```

val enter : (t -> unit tactic) -> unit tactic

```

Now we see that our tactic, all in all, should be a function that takes goal (`Proofview.Goal.t`) as input and return a tactic monad (`unit tactic`). When we have such a function, `enter` will apply it to each goal.

Having a `gl : Proofview.Goal.t` as input, we need to split it into the conclusion, hypotheses and the environment, using functions `concl`, `hyps` and `env` respectively. The most import for us is the `concl`, which we will

⁹<https://coq.github.io/doc/V8.16.0/api/coq-core/index.html>

get. It has a `EConstr.constr`, which is the most important datatype in Coq, namely the kernel term. It is used to represent expressions, which are built using a set of constructors that correspond to the different types and operations.

Our next task is to prepare the conclusion for the `egg` use. `Econstr.t` (similar to `EConstr.constr`) has a huge list of constructors, but we are not expecting to handle most of them. For example, constructors such as `Forall` or `LetIn` are not something we expect to see in a proposition about relations. From the limitations of what kind of rules we can pass to `egg` we can infer that we want to handle particular relations and various operations on them, i.e. applications of functions. Moreover, after analysing the Hahn library and the `imm` code base, we have decided to add some concrete relations as constants: An empty relation, denoted as `(fun _ _ => False)`, which means that for any two elements of the given set are related, and a full relation, denoted as `(fun _ _ => True)`. Thereby the constructors we are interested in are as follows:

```
type Econstr.constr =
  | App of 'constr * 'constr array
  | Var of Names.Id.t
  | Lambda of Names.Name.t Context.binder_annot * 'types * 'constr
  | Ind of Names.inductive * 'univs (* Inductive constructors, e.g. True or
    False *)
```

`Econstr.t` is parsed into a smaller type, which will be then transferred to `egg`. If unexpected constructors occur in the expression, exception is raised and a error is shown to the user. Data type is an s-expression over strings, with an addition of lambdas:

```
type goal_s_expr =
  | Symbol of string
  | Application of string * goal_s_expr list
  | Lambda of goal_s_expr * goal_s_expr
```

Next step is to pass an object of type `goal_s_expr` to Rust for further processing. We use the `ocaml-rs` [2] Rust library to set up communication between OCaml and Rust. Similar type as `goal_s_expr` is defined in Rust:

```
pub enum GoalSEExpr {  
    Symbol(String),  
    Application(String, LinkedList<GoalSEExpr>),  
    Lambda(Box<GoalSEExpr>, Box<GoalSEExpr>),  
}
```

4.3 Using Egg

References

- [1] Viktor Vafeiadis et al. “Hahn: a Coq library that contains a useful collection of lemmas and tactics about lists and binary relations”. In: (2018). URL: <https://github.com/vafeiadis/hahn>.
- [2] Zach Shipko et al. “ocaml-rs: OCaml extensions in Rust”. In: (2021). URL: <https://crates.io/crates/ocaml>.
- [3] Kenneth I. Appel. “The Use of the Computer in the Proof of the Four Color Theorem”. In: *Proceedings of the American Philosophical Society* 128.1 (1984), pp. 35–39. ISSN: 0003049X. URL: <http://www.jstor.org/stable/986491> (visited on 03/29/2023).
- [4] Mark Batty, Alastair F. Donaldson, and John Wickerson. “Overhauling SC atomics in C11 and OpenCL”. In: (2016). URL: <https://doi.org/10.1145/2F2837614.2837637>.
- [5] Yves Bertot and Pierre Castéran. *Interactive theorem proving and program development: Coq’Art: the calculus of inductive constructions*. Springer Science & Business Media, 2013.
- [6] Georges Gonthier. “Formal Proof—The Four- Color Theorem”. In: 2008. URL: <https://www.ams.org/notices/200811/tx081101382p.pdf>.
- [7] Michalis Kokologiannakis, Ori Lahav, and Viktor Vafeiadis. “Kater: Automating Weak Memory Model Metatheory and Consistency Checking”. In: *Proc. ACM Program. Lang.* 7.POPL (2023). DOI: [10.1145/3571212](https://doi.org/10.1145/3571212). URL: <https://doi.org/10.1145/3571212>.
- [8] Ori Lahav et al. “Repairing Sequential Consistency in C/C++11”. In: *SIGPLAN Not.* 52.6 (2017), pp. 618–632. ISSN: 0362-1340. URL: <https://doi.org/10.1145/3140587.3062352>.
- [9] David C. Luckham et al. *Stanford Pascal Verifier User Manual*. Tech. rep. Stanford, CA, USA, 1979.
- [10] Chandrakana Nandi et al. “Rewrite Rule Inference Using Equality Saturation”. In: *Proc. ACM Program. Lang.* 5.OOPSLA (2021). DOI: [10.1145/3485496](https://doi.org/10.1145/3485496). URL: <https://doi.org/10.1145/3485496>.

- [11] Greg Nelson and Derek C. Oppen. “Fast Decision Procedures Based on Congruence Closure”. In: *J. ACM* 27.2 (1980), pp. 356–364. ISSN: 0004-5411. DOI: [10.1145/322186.322198](https://doi.org/10.1145/322186.322198). URL: <https://doi.org/10.1145/322186.322198>.
- [12] Christine Paulin-Mohring. “Introduction to the Calculus of Inductive Constructions”. In: *All about Proofs, Proofs for All*. Ed. by Bruno Woltzenlogel Paleo and David Delahaye. Vol. 55. Studies in Logic (Mathematical logic and foundations). College Publications, Jan. 2015. URL: <https://hal.inria.fr/hal-01094195>.
- [13] Jean Pichon-Pharabod and Peter Sewell. “A Concurrency Semantics for Relaxed Atomics That Permits Optimisation and Avoids Thin-Air Executions”. In: *SIGPLAN Not.* 51.1 (2016), pp. 622–633. ISSN: 0362-1340. DOI: [10.1145/2914770.2837616](https://doi.org/10.1145/2914770.2837616). URL: <https://doi.org/10.1145/2914770.2837616>.
- [14] Robert Endre Tarjan. “Efficiency of a Good But Not Linear Set Union Algorithm”. In: *J. ACM* 22.2 (1975), pp. 215–225. ISSN: 0004-5411. DOI: [10.1145/321879.321884](https://doi.org/10.1145/321879.321884). URL: <https://doi.org/10.1145/321879.321884>.
- [15] Max Willsey et al. “Egg: Fast and Extensible Equality Saturation”. In: *Proc. ACM Program. Lang.* 5.POPL (2021). DOI: [10.1145/3434304](https://doi.org/10.1145/3434304). URL: <https://doi.org/10.1145/3434304>.